

Virtualization (Task Merging) Design & Implementation

-


Embedded Software Group – Delft University of Technology, The Netherlands

8 February 2017



FRAM Write

- Almost twice more expensive than SRAM write.

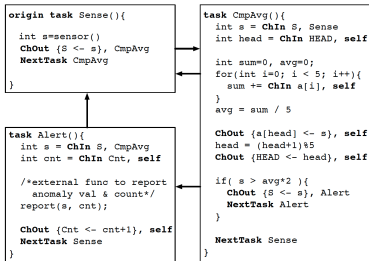


Problem Statement

- How to **postpone** commit (FRAM access) operations to reduce energy overhead?
 - The gained energy will be used for **more computation**
 - E.g. more tasks will be executed in task-based programs — in turn, the total program execution time will be reduced
 - Portability: different super-capacitor sizes will be used efficiently

Task-based Control Flow

- e.g. Chain¹
 - no checkpointing – keep track of the current task
 - pipelined channels – maintain the consistency of FRAM
 - separate inputs from outputs – idempotency




Drawbacks

- outputs are committed **at the end** of each task
 - **Commit** operation, in turn FRAM access (in particular write), is **expensive**
- **multi-versioning** of variables — memory overhead
 - increased number of FRAM operations

¹colin2016chain

Details - Virtualization of Chain -I

- If all output channels are connected to and the transition path is only to the **next task**, virtualization is **easy**!
- No need to commit Channel (1) and Channel (2) since their values are already consumed to output Channel (3).




Strategy

- Postpone commit inside the **virtualization window** using corresponding **soft-channels** in SRAM
- **Hard-commit** only the channel input of the last task
- If power-interrupted, go to the first task of the window

Details - Virtualization of Chain -II

- What if the channel connections and transition paths are **divergent**?
 - Channel (3), Channel (4) and Channel (5) **must** committed to FRAM—**hard-channels**
 - Channel (1) and Channel (2) are in SRAM—**soft-channels**
 - Channel (6) should be read from FRAM—hard-channel



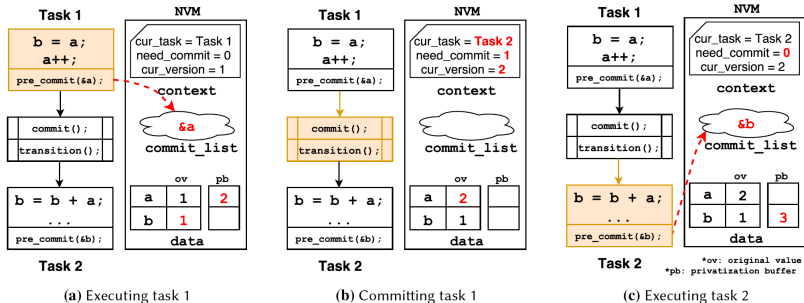
Drawbacks

- The static **task-graph** and the **dynamic** task transition path should be used at **run-time** to decide which channels need to be hard or soft

Not Feasible

- Due to loops in task graph, also multi-versioning ruins virtualization.

Case Study - Virtualization of ALPACA - I




Only variables with W-A-R dependencies are protected in ALPACA.

Challenges

- Variables without W-A-R dependencies are directly read from/written to FRAM
- Solution:** Use **SRAM buffer** and commit later—Reduced FRAM access

Case Study - Virtualization of ALPACA -II

- Merged tasks might introduce new W-A-R dependencies —at **run-time**!
- How to **dynamically** keep-track of the privatization of the W-A-R dependencies? —**too much overhead**!
 - Alpaca does not privatize x , y and z ; they are directly written to FRAM.
 - We need software implementation of Clank :)

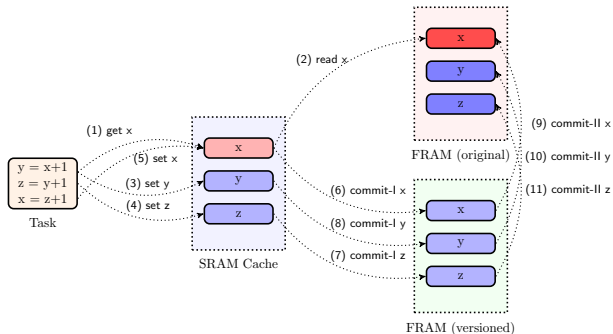


Conclusions

- We need to **eliminate** privatization so that tasks can operate on SRAM efficiently.

Virtualization with SRAM Buffer-I

- Version **task-shared** variables in FRAM
- A **fixed-sized** SRAM buffer for virtual operations
- Two-phase commit at the end of the **virtual task**.




Overhead

- get and set operations require **search** in SRAM cache
- Alpaca will have 4 FRAM write operations - We will have 6 + SRAM search overhead!

Virtualization with SRAM Buffers-II

- We did not observe any benefit, even worse!





Problem

- How to eliminate SRAM Cache search operation?
- How to minimize commit overhead?

Virtualization with DMA Support-I

- Exploit hardware that is found in almost every CPU architecture
- Copying block of words from/to SRAM-FRAM and FRAM-FRAM is very **cheap** and **scalable** with DMA!

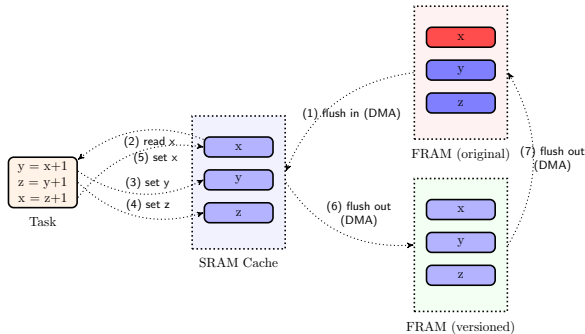


Problem

- How to eliminate SRAM Cache search operation?
- How to minimize commit overhead?

Virtualization with DMA Support-II



- Version all **task-shared** variables in FRAM
- Bring all to SRAM buffer for virtual operations
- Two-phase commit at the end of the **virtual task**.



Feasibility

- Perform all these operations using DMA!

Virtualization with DMA Support-III



Virtualization with DMA Support-IV

- We have some limitations:
 - The size of the task-shared variables, in turn the size of SRAM Cache is bounded with the **SRAM size** in the platform.
 - The number of channels is directly limited with the FRAM size in Chain
 - The size of the commit list in Alpaca, i.e. the number of variables with W-A-R dependency inside a task, is limited with the FRAM.
 - The cost of start-up and commit operations increase with **the size of the task-shared variables**
 - But this comes with the elimination of the search operations in SRAM cache and selective commit
 - DMA cost increases linearly but very slowly.
- all these limitations are paid in order to benefit from virtualization.

What is next?

- How to get rid of flush in/out all task-shared global variables?
- Development and evaluation of several task-merging algorithms, i.e. **virtualization scheduler**

The structure of the article-I

- **Introduction**

- Motivation—FRAM access, in turn commit cost, need for the virtualization
- Problem statement
- Challenges—Why it is difficult to modify state of the art solutions to support virtualization?
- Our solution and contributions—Sketch of our solution and attractive results

- **Related Work and Challenges**

- Intermittent programming with checkpointing and task-based control flow
- Stress the problem of dynamic merging of the tasks, its feasibility with the state of the art solutions, e.g. Chain, InOS, Alpaca...

- **Virtualization: Dynamic Merging of Tasks**

- Introduce the concept, requirements
- Design objectives
- Introduce the infrastructure here

The structure of the article-II

- **Virtualization Scheduler: The Algorithms**

- Introduce the task-merging strategies that work on the infrastructure in the previous section
- Analysis of the algorithms
- Maybe some simulations

- **Implementation and Evaluation**

- Introduce the implementation details
- Test-bed setup
- Experimental results
- Simulations

- **Conclusions**